

# RFC 0.1: Semantic Intent Negotiation Protocol (SINP)

Specification Document  
Version: 0.1 (Verified Draft)

December 31, 2025

## Abstract

This document specifies the Semantic Intent Negotiation Protocol (SINP), an application-layer protocol designed to replace deterministic address-routing with semantic negotiation. It defines a rigorous state machine where endpoints exchange *Intent Descriptions* and resolve execution through mutual *Confidence* scores. This version (0.1) includes explicit wire formats, cryptographic canonicalization rules, decision-theoretic thresholds, and failure modes for probabilistic interpreters (LLMs).

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
1.1	Terminology . . . . .	2
<b>2</b>	<b>Core Concepts and Formal Definitions</b>	<b>2</b>
2.1	The Message Tuple . . . . .	2
2.2	Confidence ( $\Phi$ ) . . . . .	2
<b>3</b>	<b>Wire Format Specification</b>	<b>2</b>
3.1	Request Schema . . . . .	2
3.2	Response Schema . . . . .	3
3.3	Capability Schema . . . . .	3
<b>4</b>	<b>Mathematical Model and Decision Logic</b>	<b>3</b>
4.1	Interpretation Function . . . . .	3
4.2	Confidence Derivation . . . . .	4
4.3	Decision Boundary ( $\delta$ ) . . . . .	4
<b>5</b>	<b>State Machines</b>	<b>4</b>
5.1	Server State Automaton . . . . .	4
5.2	Client State Automaton . . . . .	4
<b>6</b>	<b>Interpretation Algorithms</b>	<b>4</b>
6.1	Deterministic Scoring (Baseline) . . . . .	4
6.2	LLM Calibration Observability . . . . .	4
<b>7</b>	<b>Security Integrity</b>	<b>5</b>
7.1	Semantic Hashing (Caching) . . . . .	5
7.2	Replay Protection . . . . .	5
7.3	Signature Canonicalization . . . . .	5
<b>8</b>	<b>Appendix: Refusal Codes</b>	<b>5</b>

# 1 Introduction

SINP bridges the gap between deterministic systems and probabilistic intent (NLP/LLMs). It formalizes the exchange of *Intent*, *Context*, and *Confidence* to allow a server to clarify, propose alternatives, or refuse requests deterministically before execution.

## 1.1 Terminology

The key words "MUST", "MUST NOT", "REQUIRED", "SHOULD", and "OPTIONAL" are to be interpreted as described in RFC 2119.

# 2 Core Concepts and Formal Definitions

## 2.1 The Message Tuple

A message  $M$  in SINP is defined as a tuple:

$$M = (ID, CID, T, Sender, \Psi, \Gamma, \Phi, \Sigma) \quad (1)$$

Where:

- $ID$ : UUIDv4 message identifier.
- $CID$ : UUIDv4 conversation identifier.
- $T$ : ISO-8601 UTC Timestamp.
- $Sender$ : Identity object  $\{id, auth\}$ .
- $\Psi$ : The Intent (Request) or Interpretation (Response) text.
- $\Gamma$ : The Context object (Transcript/Summary).
- $\Phi$ : The scalar Confidence score,  $\Phi \in [0, 1]$ .
- $\Sigma$ : Cryptographic signature (Optional).

## 2.2 Confidence ( $\Phi$ )

Confidence is a scalar value representing *willingness to act*.

- $\Phi_c$  (Client): Certainty that  $\Psi$  reflects the user's desire.
- $\Phi_s$  (Server): Probability that the mapped capability  $c$  satisfies  $\Psi$  given  $\Gamma$ , adjusted for policy and safety.

# 3 Wire Format Specification

All messages MUST be serialized as JSON.

## 3.1 Request Schema

**Note:** The first message in a conversation MUST omit 'in\_response\_to'. All subsequent messages MUST include it.

```
{  
  "protocol_version": "0.1",  
  "message_id": "uuid-v4",  
  "in_response_to": "uuid-v4 (OPTIONAL for INIT)",  
  "conversation_id": "uuid-v4",  
  "timestamp": "2025-12-31T12:00:00Z",  
  "sender": { "id": "client_1", "auth_method": "token" },  
  "intent": "string",  
  "confidence": 0.85, // REQUIRED, must be > 0  
  "context": {
```

```

    "type": "transcript",
    "content": "...",
    "semantic_hash": "sha256_hex"
},
"constraints": {
    "max_cost": 10,
    "privacy": "private"
},
"signature": "base64_sig (OPTIONAL)"
}

```

Listing 1: Client Request Structure

### 3.2 Response Schema

```
{
    "message_id": "uuid-v4",
    "in_response_to": "uuid-v4",
    "conversation_id": "uuid-v4",
    "timestamp": "2025-12-31T12:00:01Z",
    "responder": {
        "id": "srv_1",
        "capabilities": ["cap_id_1:v1"]
    },
    "interpretation": {
        "text": "Mapped explicit action description",
        "confidence": 0.92
    },
    "action": "EXECUTE | CLARIFY | PROPOSE | REFUSE",
    "action_metadata": {
        // Contains 'result' if EXECUTE
        // Contains 'questions' if CLARIFY
        // Contains 'reason_code' if REFUSE
    },
    "alternatives": [ // OPTIONAL
    {
        "interpretation": "Alternative action description",
        "confidence": 0.80,
        "estimated_cost": 1.5,
        "capability_id": "cap_id_2:v1"
    }
    ],
    "confidence": 0.90
}

```

Listing 2: Server Response Structure

### 3.3 Capability Schema

Servers MUST define capabilities using stable identifiers. Clients MAY query these via a DISCOVER intent or out-of-band registry.

```
{
    "id": "fetch_profile:v1",
    "description": "Returns user profile data.",
    "inputs": ["user_id"],
    "privacy_level": "pii_sensitive",
    "cost_units": 1.0
}
```

## 4 Mathematical Model and Decision Logic

### 4.1 Interpretation Function

The server implements an interpretation function  $f$ :

$$f(\Psi_{req}, \Gamma) \rightarrow (\hat{\Psi}, c, \rho) \quad (2)$$

Where  $\hat{\Psi}$  is the server's interpretation,  $c$  is the capability, and  $\rho$  is the raw model probability.

## 4.2 Confidence Derivation

Let  $R(c)$  be the reliability factor of capability  $c$ ,  $A(res)$  be resource availability, and  $P(pol) \in \{0, 1\}$  be the policy check.

$$\Phi_s = \min(1, \rho \cdot R(c) \cdot A(res)) \cdot P(pol) \quad (3)$$

## 4.3 Decision Boundary ( $\delta$ )

The server determines Action  $A$  using thresholds  $\tau_{exec}$  (default 0.85),  $\tau_{clarify}$  (default 0.50), and  $\tau_{accept}$  (default 0.50).

$$A = \delta(\Phi_s, \Phi_c) = \begin{cases} \text{EXECUTE} & \text{if } \Phi_s \geq \tau_{exec} \wedge \Phi_c \geq \tau_{accept} \\ \text{REFUSE} & \text{if } P(pol) = 0 \vee \text{Malformed} \\ \text{PROPOSE} & \text{if } \exists c' \neq c : \Phi_s(c') > \Phi_s(c) \\ \text{CLARIFY} & \text{if } \Phi_s < \tau_{exec} \end{cases} \quad (4)$$

# 5 State Machines

## 5.1 Server State Automaton

1. **RECEIVED**: Validate signature (JCS) and schema. Check Replay ( $|T_{now} - T_{msg}| \leq 5s$ ).
2. **INTERPRETING**: Run  $f(\Psi, \Gamma)$ . Compute  $\Phi_s$ .
3. **DECIDING**: Apply  $\delta(\Phi_s, \Phi_c)$ .
  - EXECUTE → Perform Action → State **DONE**.
  - CLARIFY → Generate Questions → State **NEGOTIATING**.
  - PROPOSE → Generate Alternatives → State **NEGOTIATING**.

## 5.2 Client State Automaton

- **INIT**: Send Request (no `in_response_to`).
- **REFINING**: Wait for Response.
  - On CLARIFY: User input → Update Context → Send Request.
  - On PROPOSE:
    - \* If Accepted: Send new Request referencing proposal ID.
    - \* If Rejected: Send new Request with modified intent.
  - On EXECUTE: Terminate (**SATISFIED**).

# 6 Interpretation Algorithms

## 6.1 Deterministic Scoring (Baseline)

For capability keywords  $K_c$  and intent vector  $V_{int}$ :

$$\text{Score}(c) = \frac{\sum_{k \in K_c} \mathbb{I}(k \in V_{int})}{|K_c|} \cdot w_{match} \quad (5)$$

## 6.2 LLM Calibration Observability

If using LLMs, raw confidence  $x$  MUST be calibrated (e.g., Platt Scaling). Servers SHOULD publish monitoring metrics, specifically the **Brier Score** (BS) to ensure claimed confidence matches empirical accuracy:

$$BS = \frac{1}{N} \sum_{t=1}^N (f_t - o_t)^2 \quad (6)$$

Where  $f_t$  is the forecast probability ( $\Phi_s$ ) and  $o_t$  is the outcome (1 if user accepted/successful, 0 otherwise).

## 7 Security Integrity

### 7.1 Semantic Hashing (Caching)

The semantic hash is used for caching interpretations. It MUST exclude timestamps to ensure identical intents hit the cache.

$$H_{sem} = \text{SHA256}(\text{normalize}(\Psi) \parallel \text{normalize}(\Gamma)) \quad (7)$$

### 7.2 Replay Protection

Replay attacks are mitigated by checking the tuple  $(ID, CID, T)$ . Servers MUST reject messages where  $|T_{now} - T_{sender}| > 5000ms$  (configurable).

### 7.3 Signature Canonicalization

If `signature` is used, the JSON payload MUST be canonicalized using **JCS (RFC 8785)**.

1. Remove `signature` field.
2. Canonicalize remaining JSON.
3. Sign using sender's private key.
4. Append signature to original JSON.

## 8 Appendix: Refusal Codes

Common `action_metadata.reason_code` values:

- `malformed_context`: Semantic hash mismatch or invalid structure.
- `privacyViolation`: Request requires PII but privacy constraints forbid.
- `capabilityMissing`: No capability matches intent with  $\Phi > 0.2$ .
- `policyViolation`: Intent understood but forbidden by server rules.